

Kubernetes Clustering

In this section we will cover

- Cluster Architecture
- High Availability
- Multi Cluster/Region

Architecture

Master



etcd

api-server
controller-manager
scheduler
etcd

Node



kubelet
kube-proxy

Main
page

Main
page

Node



kubelet
kube-proxy

Main
page

Redis

Node



kubelet
kube-proxy

Redis

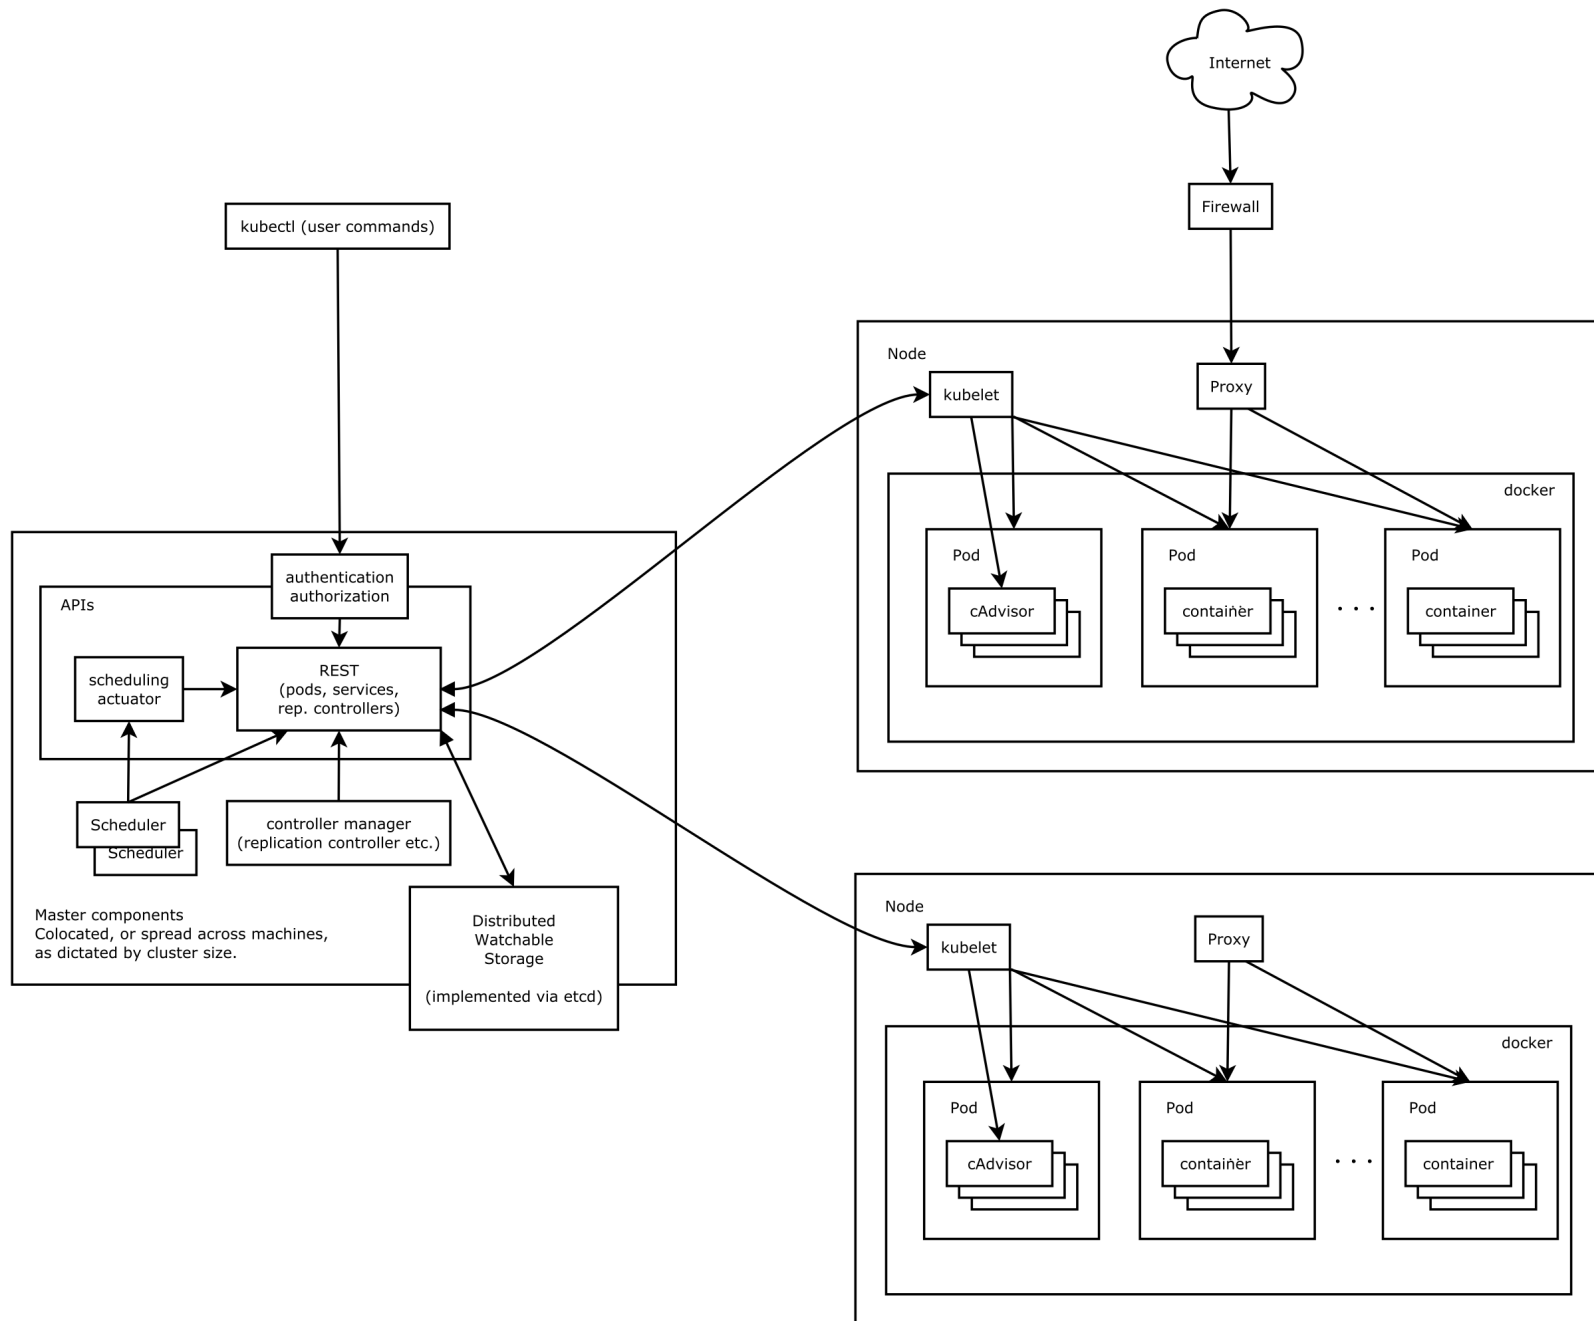
Rese
rvatio
ns

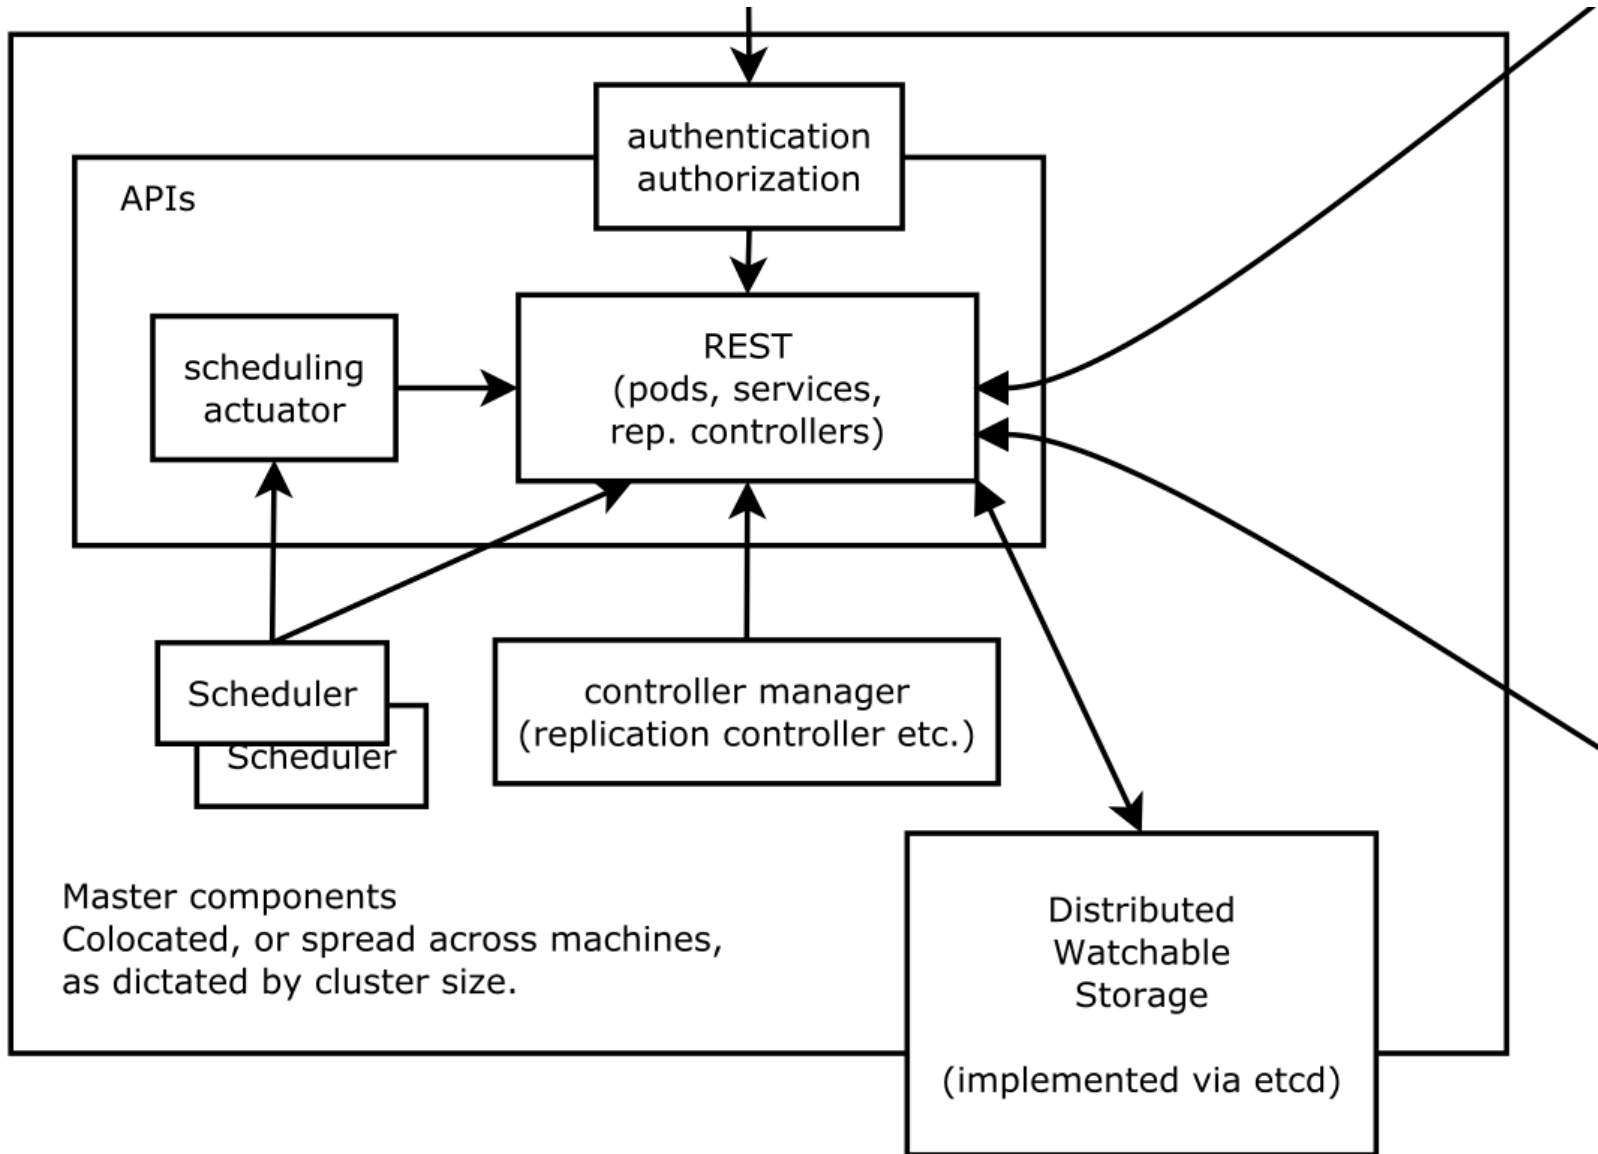
Node

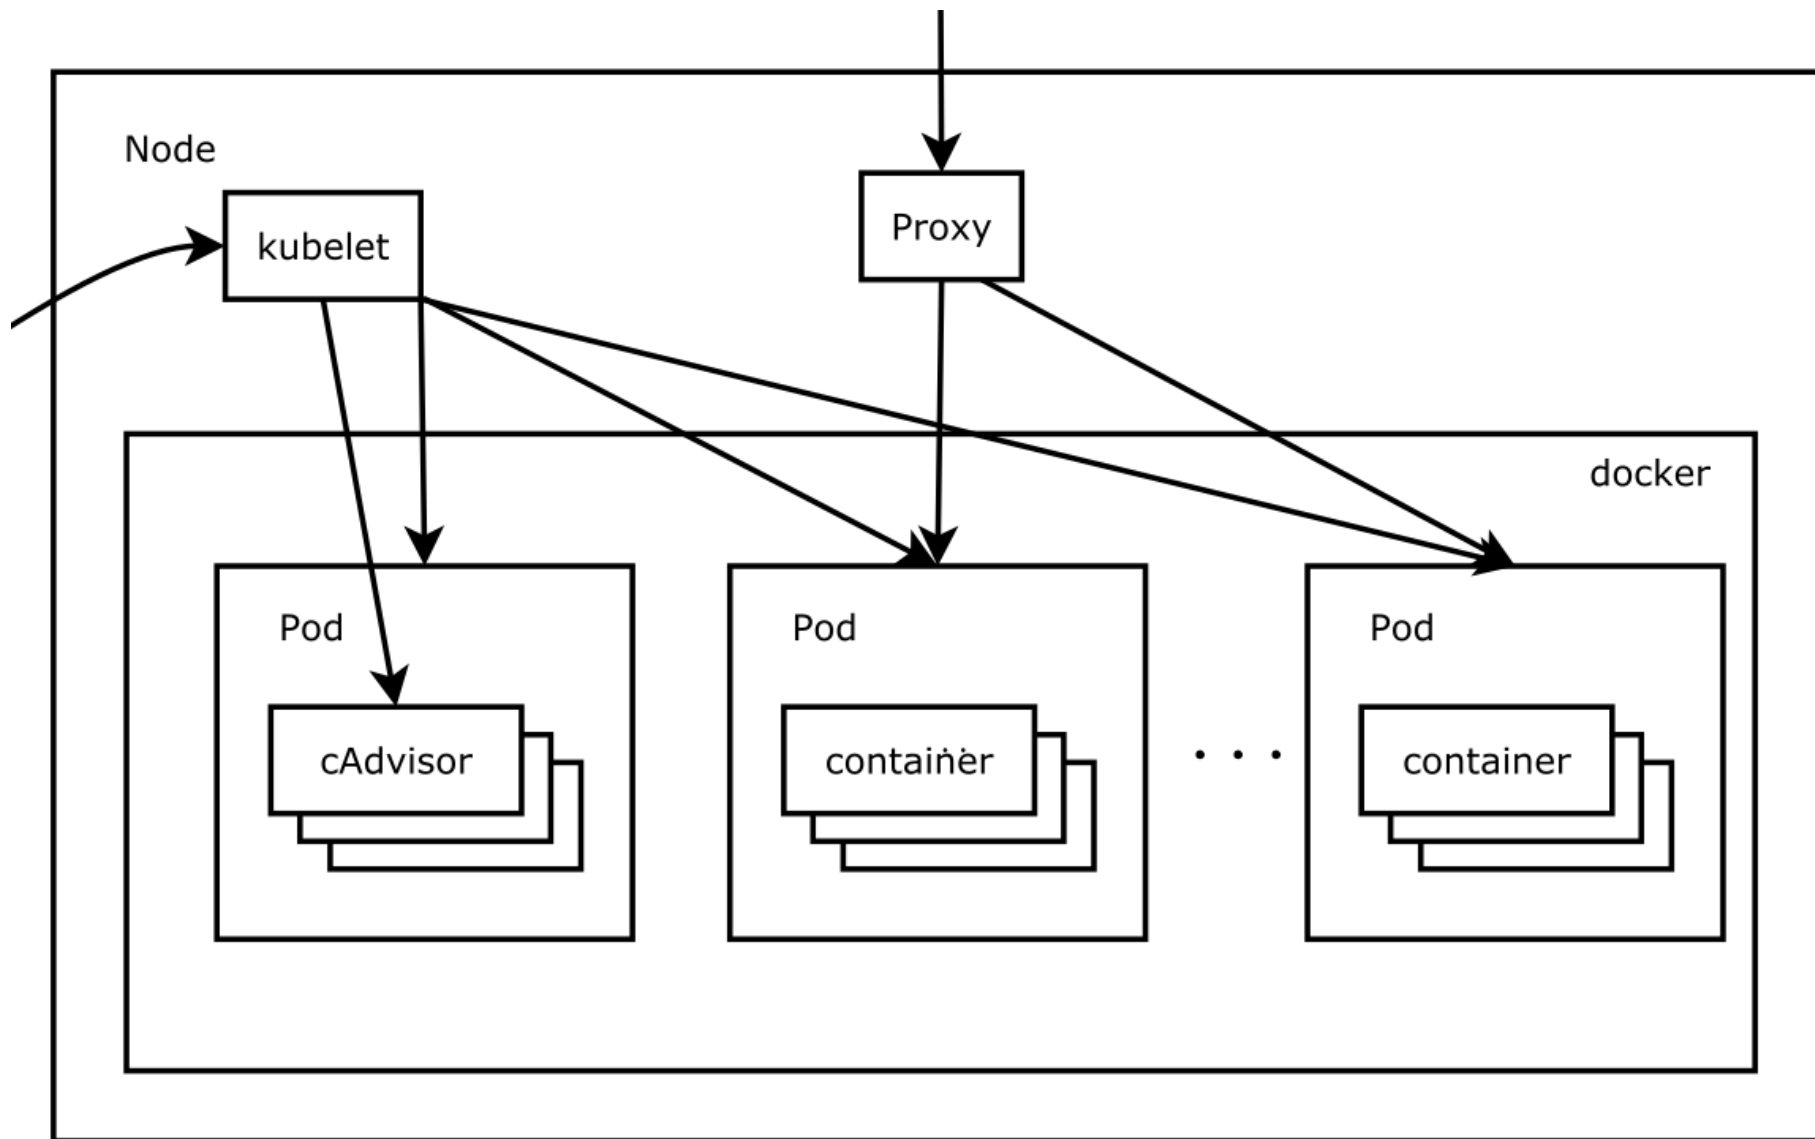


kubelet
kube-proxy

Redis







Master

- api-server
 - Provides outbound Kubernetes REST API
 - Validates requests
 - Saves cluster state to etcd

Master

- controller-manager
 - Runs "control loops"
 - Regulates the state of the system
 - Watches cluster state through the api-server
 - Changes current state towards the desired state
 - e.g. checks correct number of pods running

Master

- Scheduler
 - Selects node on which to run a pod

Master

- etcd
 - Distributed, consistent key-value store for shared configuration and service discovery

Node

- kubelet
 - Agent that runs on each node
 - Takes a set of PodSpecs from API server
 - Starts containers to fulfill specs
 - Exposes monitoring data

Node

- kube-proxy
 - Implements service endpoints (virtual IPs)
 - IPTables

**High Availability
And Multi Region**

Highly Available

Why do we want our systems to be Highly Available?

"Everything fails, all the time"

-Werner Vogels

What does HA even mean?

High availability is a characteristic of a system, which aims to ensure an agreed level of operational performance, usually uptime, for a higher than normal period.

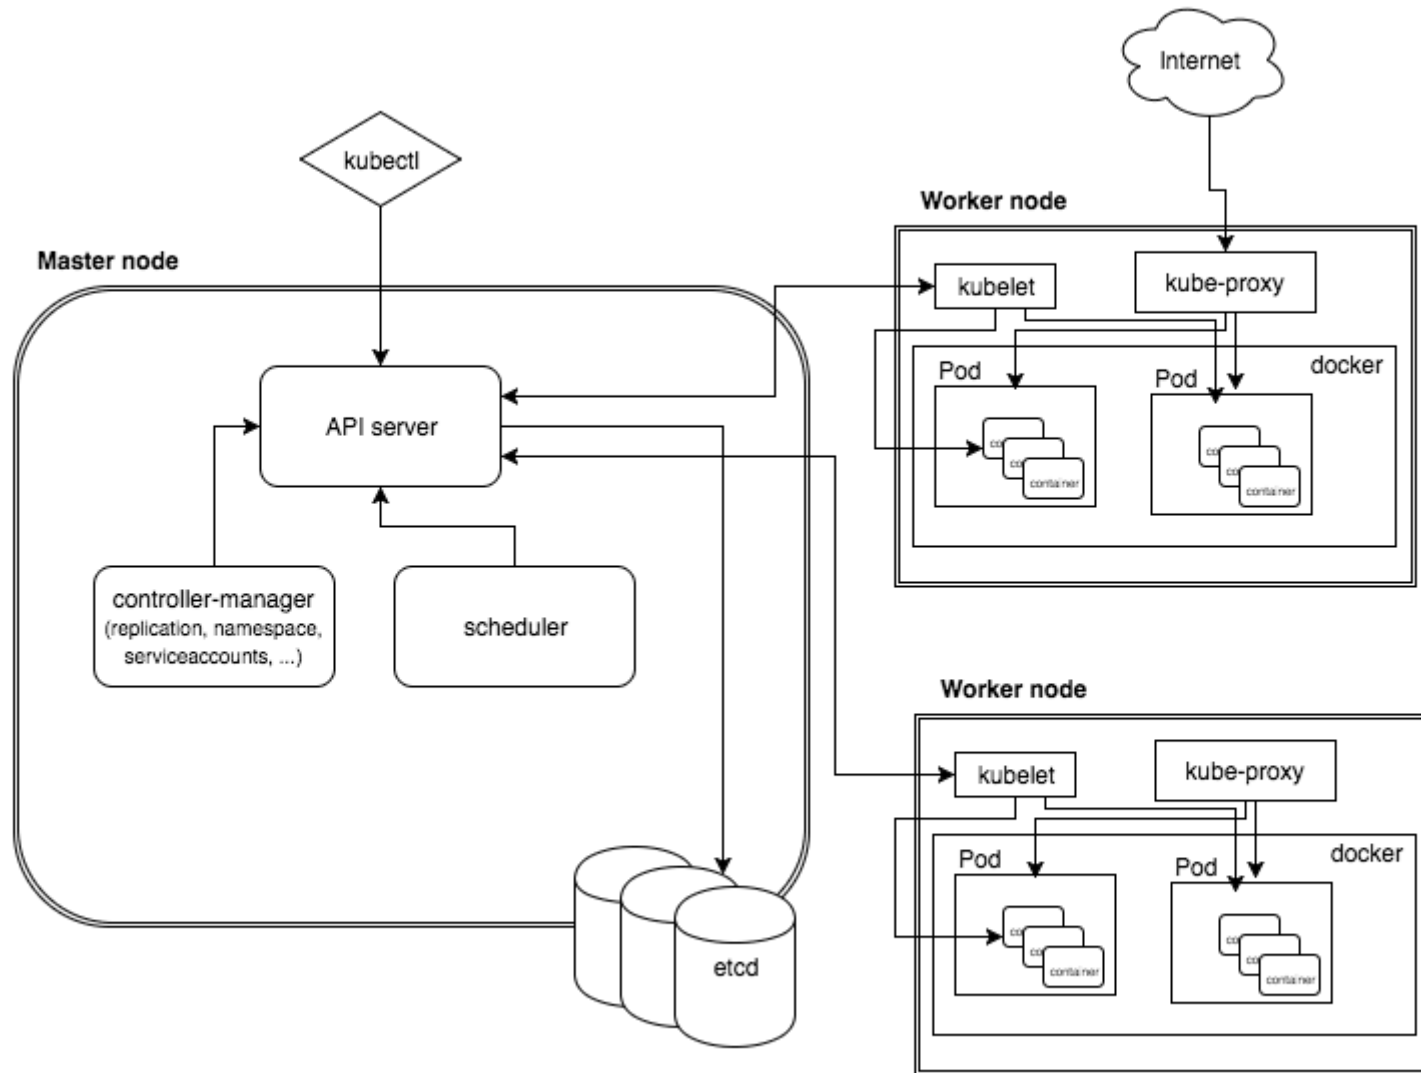
-Wikipedia

(This sounds a bit vague)

When we talk about systems which are Highly Available we mean that there should be no **single point of failure**

What are the potential single points of failure in a Kubernetes cluster?

Kubernetes Architecture



HA Kubernetes

Master

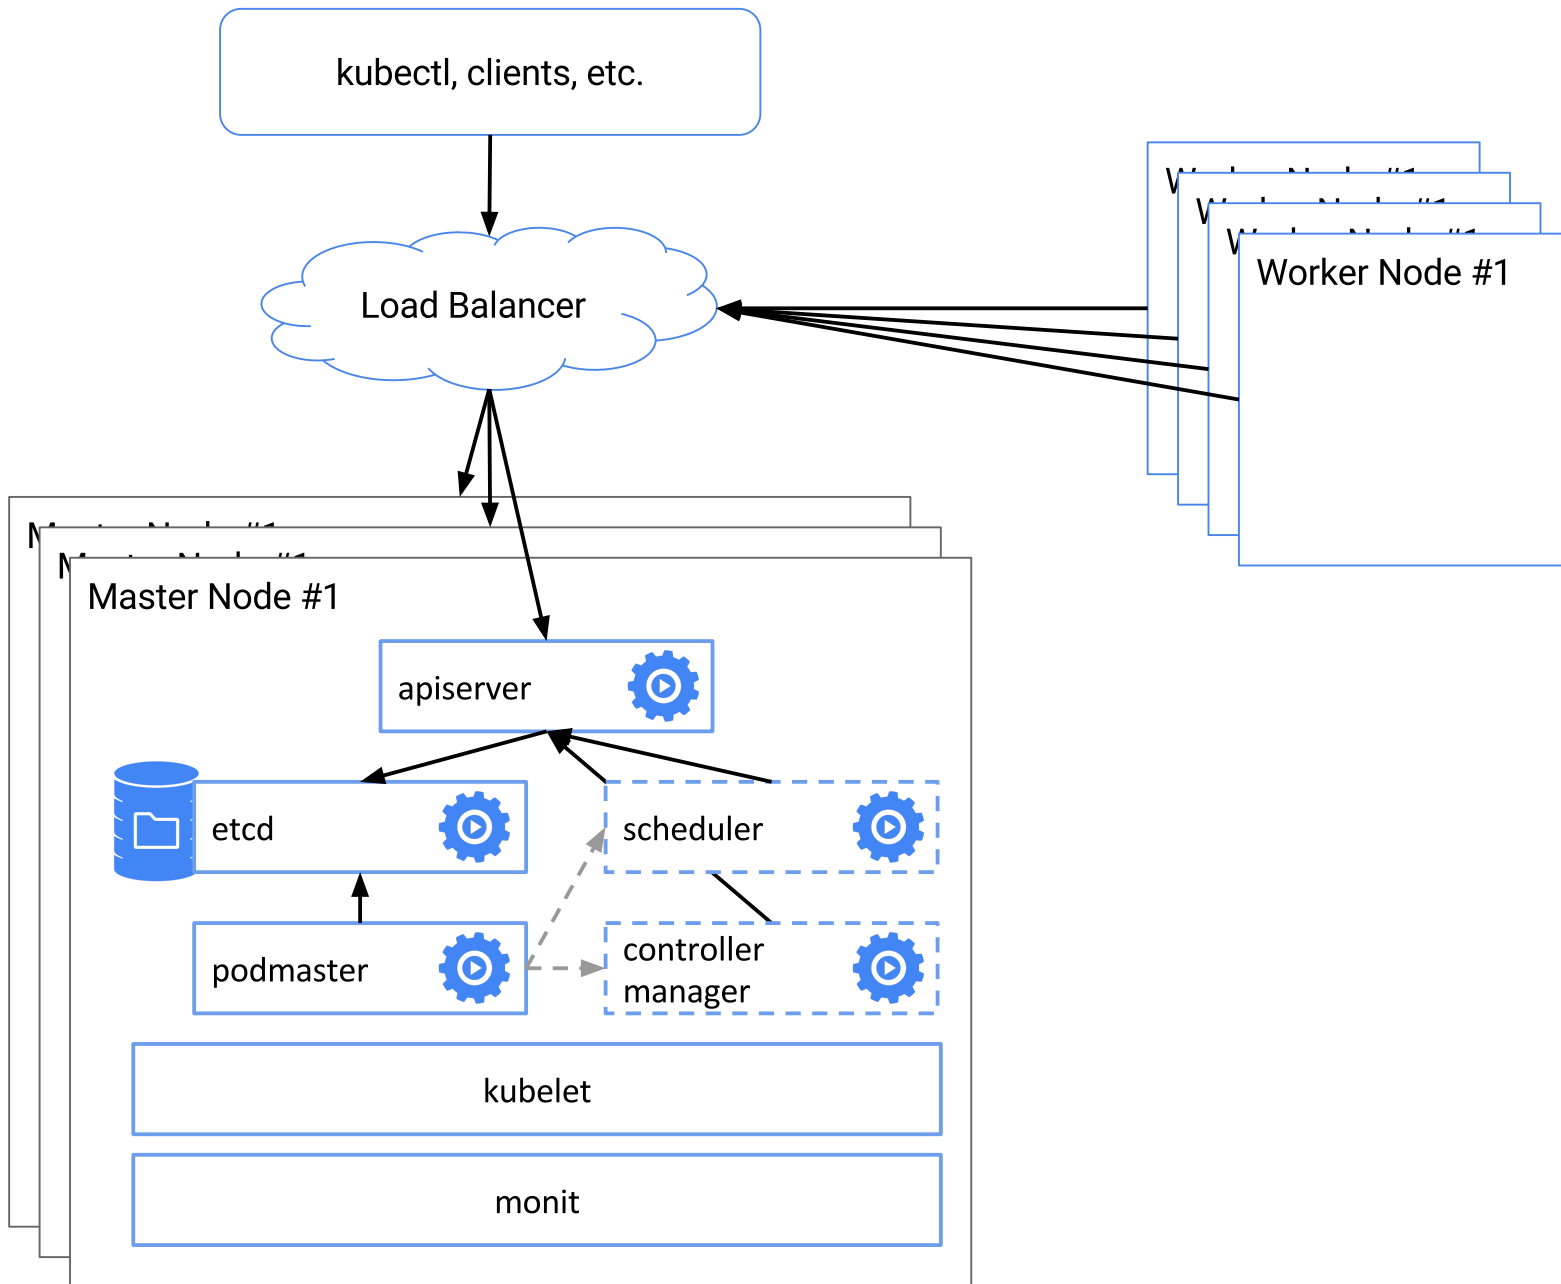
- API Server
- Controller Manager*
- Scheduler*

*Will perform leader election (given the `--leader-elect` flag)

HA Kubernetes

Etcd

- Etcd maintains the state of our cluster. This is crucial to maintain high availability.
- We can accomplish by creating a (minimum) 3 node Etcd cluster.
- We can also use persistent storage (e.g. on a cloud provider) to ensure no data is lost.



To co-locate Masters and Etcd nodes?

This is a trade-off between managing/paying for more instances vs isolation.

Load Balancing the API Server

- Need to handle master failure
- Kube-proxies need to point to API Server
- kubectl (or any other integration as well)

Common HA Setup on the Cloud
Spread across 3 regions

How to setup a Highly Available Kubernetes cluster?

Easy :)

```
$ gcloud container clusters create my-first-cluster
```

Depends on your Infrastructure

- Google Kubernetes Engine
- Amazon EKS
- Azure Container Service
- OpenShift Origin
- Giant Swarm
- Kubermatic
- Tectonic by CoreOS
- IBM Cloud Container Service
- Kubespray
- Kops
- Kube-Up
- Kubeadm
- Kubicorn

Depends on your Needs

- Budget (time and money)
- Are you running in the cloud or onprem?
- Do you have dedicated infra/ops team?
- What are your security requirements?
- Do you like to do things the hard way :)?

Role-based Access Control (RBAC)

A bit about Roles

Role-based access control (RBAC) is a common approach to managing users' access to resources or operations.

Permissions specify exactly which resources and actions can be accessed.

The basic principle is: instead of separately managing the permissions of each user, permissions are given to roles, which are then assigned to users, or better - groups of users.

Roles Bundle Permissions

- Managing permissions per user can be a tedious task when many users are involved.
- As users are added to the system, maintaining user permissions becomes harder and more prone to errors.
- Incorrect assignment of permissions can block users' access to required systems, or worse - allow unauthorized users to access restricted areas or perform risky operations.

- A regular user can only perform a limited number of actions (e.g. get, watch, list).
- A closer look into these user actions can reveal that some actions tend to go together e.g. checking logs.
- Once roles are identified and assigned to each user, permissions can then be assigned to roles, instead of users.

Managing the permissions of a small number of roles
is a much easier task.

Basic concepts

Rule: grants permission

- Applies to resource types
- Grants verbs (create, edit, view, delete)
- (Cluster)Role
 - Cluster wide / within a namespace
 - List of rules
- (Cluster)RoleBinding
 - Connects (Cluster)Role to User
 - Both human & service account

API overview

The RBAC API declares four top-level types which will be covered in this section:

- Role
- ClusterRole
- RoleBinding
- ClusterRoleBinding

Role

A `Role` contains rules that represent a set of permissions. Permissions are additive (there are no “deny” rules).

A `Role` can be defined within a namespace, or cluster-wide (`Role` vs `ClusterRole`)

Here's an example Role in the "default" namespace that can be used to grant read access to pods:

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""] # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

Cluster Role

- A `ClusterRole` can be used to grant the same permissions as a `Role`, but because they are cluster-scoped, they can also be used to grant access to:
 - cluster-scoped resources (like nodes)
 - namespaced resources (like pods) across all namespaces

Cluster Role

The following ClusterRole can be used to grant read access to secrets in any particular namespace, or across all namespaces

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  # "namespace" omitted since ClusterRoles are not namespaced
  name: secret-reader
rules:
- apiGroups: ["" ]
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]
```

RoleBinding

A role binding grants the permissions defined in a role to a user

Permissions can be granted within a namespace with a `RoleBinding`, or cluster-wide with a `ClusterRoleBinding`.

A RoleBinding may reference a Role in the same namespace. The following RoleBinding grants the “pod-reader” role to the user “jane” within the “default” namespace.

```
# This role binding allows "jane" to read pods in the "default" namespace
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: User
  name: jane
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

In this example, even though the following RoleBinding refers to a ClusterRole, **dave** will only be able read secrets in the **development** namespace (the namespace of the RoleBinding).

```
# This role binding allows "dave" to read secrets in the "development" namespace
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: read-secrets
  namespace: development # This only grants permissions within the "development" namespace
subjects:
- kind: User
  name: dave
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: secret-reader
  apiGroup: rbac.authorization.k8s.io
```

Cluster Role Binding

A `ClusterRoleBinding` may be used to grant permission at the cluster level and in all namespaces. The following `ClusterRoleBinding` allows any user in the group “manager” to read secrets in any namespace.

```
# This cluster role binding allows anyone in the "manager" group to read secrets in any namespace
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: read-secrets-global
subjects:
- kind: Group
  name: manager
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: secret-reader
  apiGroup: rbac.authorization.k8s.io
```

Next up Monitoring...